

# TP2 - Foundations of Parallel Computing

Oussama Laaroussi

February 4, 2026

## Exercise 1: Loop Optimizations

We analyzed the impact of manual loop unrolling on execution time for summation loops using `int` and `float` types, compiled with `-O0` and `-O2`.

### Code Analysis

The loop was unrolled with factors  $U = 1, 2, 4, \dots, 32$ .

```
1 // Example U=4
2 for (int i = 0; i < N; i+=4)
3     sum += a[i] + a[i+1] + a[i+2] + a[i+3];
```

Listing 1: Manual Unrolling Structure

## Results




Figure 1: Execution Time vs Unrolling Factor

- **Float (O0):** Manual unrolling significantly reduces time (from  $\approx 29\text{ms}$  to  $\approx 9\text{ms}$  at  $U = 4$ ). This is due to reduced loop overhead (fewer comparisons/increments) and better instruction scheduling potential.
- **Int (O0):** Similar trend, with optimized unrolling reaching  $\approx 7.5\text{ms}$ .
- **Comparison with O2:**
  - For **int**, the compiler optimization (**-O2**) achieves  $\approx 2.7\text{ms}$ , beating the best manual unrolling.
  - For **float**, curiously, the manual unrolling at **-O0** (9ms) performed slightly better than the baseline **-O2** (13.5ms). This might indicate that the compiler prioritized strict floating-point associativity over aggressive vectorization in this specific simple case.

## Exercise 2: Instruction Scheduling

We improved performance by manually breaking dependency chains to allow the CPU pipelined execution units to work in parallel.

### Optimization Technique

The original code accumulated results into variables **x** and **y** sequentially. In the optimized version, we compute a common term **res = a \* b** and then accumulate into **x** and **y** in a manner that exposes more independent operations for the CPU pipeline.

```

1 int main() {
2     double a = 1.1, b = 1.2;
3     double x = 0.0, y = 0.0;
4     clock_t start, end;
5
6     start = clock();
7     double res = a * b; // Calculate once
8
9     // Loop unrolling implicitly handled or simple iteration
10    for (int i = 0; i < N/4; i++) {
11        // Breaking dependencies:
12        // Using 'res' allows independent adds if the compiler
13        // or hardware can pipeline it.
14        x = res + res + res + res + x;
15        y = res + res + res + res + y;
16    }
17    end = clock();
18    // ...
19 }
```

Listing 2: Manually Optimized Version (ex2\_manually\_optimized.c)

## Performance Comparison

- **Original -O0:** 0.231s
- **Original -O2:** 0.103s
- **Manually Optimized:** 0.066s

The manual optimization beats the compiler optimization (-O2) because it fundamentally changes the algorithm to reduce the number of multiplications and exposes 4 independent additions per loop iteration.

## Exercise 3: Mixed Workload (Scaling)

We profiled a program containing sequential (noise generation, reduction) and parallelizable (initialization, addition) parts.

### Profiling Data (Callgrind)

For  $N = 10^8$ :

- `compute_addition` (Parallel): 33.85%
- `add_noise` (Sequential): 27.69%
- `init_b` (Parallel): 20.00%
- `reduction` (Sequential): 18.46%

**Sequential Fraction ( $f_s$ ):**  $\approx 27.7\% + 18.5\% = 46.2\%$ .

We analyzed a program where the sequential fraction  $f_s$  remains constant regardless of problem size.

### Data Analysis ( $f_s \approx 46\%$ )

Since all parts of the algorithm (initialization, noise, addition, reduction) scale linearly as  $O(N)$ , the ratio between parallel and sequential work remains constant at  $f_s \approx 0.46$ .

## Scaling Analysis



Figure 2: Strong Scaling (Amdahl's Law) for Ex3




Figure 3: Weak Scaling (Gustafson's Law) for Ex3

Due to  $f_s \approx 46\%$ , Strong Scaling (Fig 2) is extremely limited ( $S_{max} \approx 2.17$ ). However, Weak Scaling (Fig 3) shows linear growth but with a very shallow slope (slope =  $1 - f_s \approx 0.54$ ). This means even if we scale the problem size with processors, we only get about 54% efficiency.

## Exercise 4: Effect of Problem Size on Parallelism

In this exercise, we observed a dramatic shift in parallel efficiency as the problem size  $N$  increased.

### Data Analysis

$f_s$  drops dramatically as  $N$  increases ( $36.6\% \rightarrow 0.35\%$ ), because the parallel part is  $O(N^3)$  while sequential parts are  $O(N)$ .

### Scaling Analysis




Figure 4: Strong Scaling (Amdahl) for Different  $N$



Figure 5: Weak Scaling (Gustafson) for Different  $N$

For large  $N = 512$ , both Strong Scaling (Fig 4) and Weak Scaling (Fig 5) are near-ideal. The Weak Scaling plot for  $N = 512$  overlaps almost perfectly with the ideal linear line, demonstrating that for algorithmically complex tasks, we can maintain high efficiency at scale.